

## Project 1: Navigation

This is for the vectorized input problem.

### Learning Algorithm

I have taken the previous exercise as a skeleton for this problem.

You have the same structure a dqn\_agent.py file that contains the learn and acting part. And a model.py file that contains the Deep Neural Network, use to evaluate the Q function.

This implementation is allowing Deep Reinforcement learning with 3 optional features

- Double DQN
- Prioritized Experience Replay
- Dueling

You can activate and deactivate using the Boolean in the hyperparameters cell.

I have also added tools, in the ./tools folder, like the prioritized replay buffer. To manage performance of big buffer, I had to implement Binary Tree sampling like proposed in the DeepMind paper.

Hyperparameters for the DQN:

All hyperparameters are manage from this cell, and are passed a argument dictionary to the relevant function.

```
In [7]: from dqn_agent import Agent

hyperparameters = {
    'dqn': {
        'n_episodes': 1300,          # Number of episode of self training
        'max_t': 1000,              # Max value of steps per episode (should be more than 200)
        'epsilon': { 'name': 'ε', 'profil': 'geometric', 'init': 1., 'bound': 0.01, 'steps': 300 }
    },
    'agent': {
        'GPU': True,                # Use GPU (Bool)
        'state_size': 37,           # State Size, here 37
        'action_size': 4,           # Action space, here 4
        'seed': 0,                  # Random generator seed
        'DDQN': True,               # Use Double DQN (Bool)
        'Prioritised_replay': False, # Use Prioritised Experience Replay (Bool)
        'Dueling': True,           # Use Dueling model (Bool)
        'BUFFER_SIZE': int(1e5),    # replay buffer size
        'BATCH_SIZE': 64,           # minibatch size
        'GAMMA': 0.99,              # discount factor
        'LR': 5e-2,                 # Learning rate
        'LEARN_EVERY': 4,           # how often to update the network
        'UPDATE_EVERY': 20,         # how often to update the network
        'alpha': { 'name': 'α', 'profil': 'constant', 'init': 0.7 },
        'TD_Error_clip': [0., 1.], #float('inf')], # Clipping of TD Error range in PER selection
        'beta': { 'name': 'β', 'profil': 'linear', 'init': 0.6, 'bound': 1., 'steps': 200*75 },
        'tau': { 'name': 'τ', 'profil': 'constant', 'init': 5e-2 },
    }
}

agent = Agent(**hyperparameters['agent'])
```


After some long trial and error, I have settled on those parameters:

- GPU active for performance.
- Double DQN active,
- Prioritized replay inactive, I have better performance without, I tried introducing clipping of TD Error, changed Alpha and Beta, etc... But still better without.
- Dueling active
- Replay buffer at 10000 like previously
- Batch size for learning at 64
- Gamma at 0.99, very standard, just not to loop.
- Learning Rate 0.05
- Trigger learning every 4 episodes
- And change the target network every 20 episodes

- Alpha, TD\_Error clipping and beta not used as PER inactive
- Tau at 0.05 as the rate to move from current to target network.

For the Neural Network

I have used a simple structure for the vectorized problem,





## Plot of Rewards

When using the “Navigation.ipynb” you have the following result in 2 plots.

1. score and average score over the last 100 episodes per episode of training.
2. Loss function and average loss over the last 100 episodes per episode of training.


With Dueling and DDQN activate



The average score over 100 episodes has reached 13 on episodes 424, as the instructor way of counting, this implies that the score of 13 was reached 100 episodes before. So 324

Episode 100	in 1.48s.	Score 3	Average Score: 0.14	Loss: 6.5e-03	LR: 5.0e-02	Paramaters $\epsilon=2.2e-01$	$\gamma=9.9e-01$	$\tau=5.0e-02$
Episode 200	in 1.47s.	Score 11	Average Score: 5.26	Loss: 1.5e-02	LR: 5.0e-02	Paramaters $\epsilon=4.6e-02$	$\gamma=9.9e-01$	$\tau=5.0e-02$
Episode 300	in 1.44s.	Score 7	Average Score: 8.98	Loss: 2.5e-02	LR: 5.0e-02	Paramaters $\epsilon=1.0e-02$	$\gamma=9.9e-01$	$\tau=5.0e-02$
Episode 400	in 1.46s.	Score 11	Average Score: 11.97	Loss: 3.9e-02	LR: 5.0e-02	Paramaters $\epsilon=1.0e-02$	$\gamma=9.9e-01$	$\tau=5.0e-02$
Episode 424	in 1.51s.	Score 22	Average Score: 13.07	Loss: 3.2e-02	LR: 5.0e-02	Paramaters $\epsilon=1.0e-02$	$\gamma=9.9e-01$	$\tau=5.0e-02$
Environment solved in 324 episodes! Average Score: 13.07								
Episode 500	in 1.52s.	Score 13	Average Score: 13.09	Loss: 4.0e-02	LR: 5.0e-02	Paramaters $\epsilon=1.0e-02$	$\gamma=9.9e-01$	$\tau=5.0e-02$
Episode 588	in 1.49s.	Score 23	Average Score: 14.10	Loss: 4.6e-02	LR: 5.0e-02	Paramaters $\epsilon=1.0e-02$	$\gamma=9.9e-01$	$\tau=5.0e-02$
Environment solved in 488 episodes! Average Score: 14.10								
Episode 600	in 1.47s.	Score 18	Average Score: 14.45	Loss: 5.2e-02	LR: 5.0e-02	Paramaters $\epsilon=1.0e-02$	$\gamma=9.9e-01$	$\tau=5.0e-02$
Episode 650	in 1.50s.	Score 20	Average Score: 15.03	Loss: 6.4e-02	LR: 5.0e-02	Paramaters $\epsilon=1.0e-02$	$\gamma=9.9e-01$	$\tau=5.0e-02$
Environment solved in 550 episodes! Average Score: 15.03								
Episode 700	in 1.47s.	Score 15	Average Score: 15.00	Loss: 7.7e-02	LR: 5.0e-02	Paramaters $\epsilon=1.0e-02$	$\gamma=9.9e-01$	$\tau=5.0e-02$
Episode 800	in 1.50s.	Score 20	Average Score: 14.91	Loss: 6.6e-02	LR: 5.0e-02	Paramaters $\epsilon=1.0e-02$	$\gamma=9.9e-01$	$\tau=5.0e-02$
Episode 900	in 1.48s.	Score 12	Average Score: 15.59	Loss: 6.5e-02	LR: 5.0e-03	Paramaters $\epsilon=1.0e-02$	$\gamma=9.9e-01$	$\tau=5.0e-02$
Episode 953	in 1.51s.	Score 19	Average Score: 16.03	Loss: 6.7e-02	LR: 5.0e-03	Paramaters $\epsilon=1.0e-02$	$\gamma=9.9e-01$	$\tau=5.0e-02$
Environment solved in 853 episodes! Average Score: 16.03								
Episode 1000	in 1.49s.	Score 19	Average Score: 16.20	Loss: 9.7e-02	LR: 5.0e-03	Paramaters $\epsilon=1.0e-02$	$\gamma=9.9e-01$	$\tau=5.0e-02$

Without any extra feature, just DQN



The average score over 100 episodes has reached 13 on episodes 450, as the instructor way of counting, this implies that the score of 13 was reached 100 episodes before. So 350.

```

Episode 100 in 1.25s. Score 0 Average Score: 0.21 Loss: 8.8e-03 LR: 5.0e-02 Paramaters  $\epsilon=2.2e-01$   $\gamma=9.9e-01$   $\tau=5.0e-02$ 
Episode 200 in 1.29s. Score 9 Average Score: 6.19 Loss: 2.1e-02 LR: 5.0e-02 Paramaters  $\epsilon=4.6e-02$   $\gamma=9.9e-01$   $\tau=5.0e-02$ 
Episode 300 in 1.22s. Score 4 Average Score: 7.66 Loss: 2.7e-02 LR: 5.0e-02 Paramaters  $\epsilon=1.0e-02$   $\gamma=9.9e-01$   $\tau=5.0e-02$ 
Episode 400 in 1.23s. Score 10 Average Score: 11.86 Loss: 4.0e-02 LR: 5.0e-02 Paramaters  $\epsilon=1.0e-02$   $\gamma=9.9e-01$   $\tau=5.0e-02$ 
Episode 450 in 1.26s. Score 20 Average Score: 13.02 Loss: 5.2e-02 LR: 5.0e-02 Paramaters  $\epsilon=1.0e-02$   $\gamma=9.9e-01$   $\tau=5.0e-02$ 
Environment solved in 350 episodes! Average Score: 13.02
Episode 500 in 1.24s. Score 14 Average Score: 12.60 Loss: 4.0e-02 LR: 5.0e-02 Paramaters  $\epsilon=1.0e-02$   $\gamma=9.9e-01$   $\tau=5.0e-02$ 
Episode 600 in 1.28s. Score 15 Average Score: 12.72 Loss: 5.5e-02 LR: 5.0e-02 Paramaters  $\epsilon=1.0e-02$   $\gamma=9.9e-01$   $\tau=5.0e-02$ 
Episode 700 in 1.30s. Score 13 Average Score: 13.77 Loss: 7.3e-02 LR: 5.0e-02 Paramaters  $\epsilon=1.0e-02$   $\gamma=9.9e-01$   $\tau=5.0e-02$ 
Episode 705 in 1.29s. Score 12 Average Score: 14.01 Loss: 7.6e-02 LR: 5.0e-02 Paramaters  $\epsilon=1.0e-02$   $\gamma=9.9e-01$   $\tau=5.0e-02$ 
Environment solved in 605 episodes! Average Score: 14.01
Episode 800 in 1.25s. Score 14 Average Score: 14.40 Loss: 6.0e-02 LR: 5.0e-02 Paramaters  $\epsilon=1.0e-02$   $\gamma=9.9e-01$   $\tau=5.0e-02$ 
Episode 844 in 1.26s. Score 19 Average Score: 15.03 Loss: 6.2e-02 LR: 5.0e-03 Paramaters  $\epsilon=1.0e-02$   $\gamma=9.9e-01$   $\tau=5.0e-02$ 
Environment solved in 744 episodes! Average Score: 15.03
Episode 900 in 1.28s. Score 21 Average Score: 15.15 Loss: 6.1e-02 LR: 5.0e-03 Paramaters  $\epsilon=1.0e-02$   $\gamma=9.9e-01$   $\tau=5.0e-02$ 
Episode 1000 in 1.24s. Score 8 Average Score: 15.75 Loss: 1.0e-01 LR: 5.0e-03 Paramaters  $\epsilon=1.0e-02$   $\gamma=9.9e-01$   $\tau=5.0e-02$ 



```

My conclusion is that without extra features it takes more time to learn and you get a lower best score. Probably due to the difficulty of randomly walking into a yellow banana, it seems to take 100 episodes to understand that grabbing yellow bananas is better than wandering around. But even more than the extra features, I fear that hyperparameters right picking is also very influential on the end result.

## Ideas for Future Work

On the vector problem:

- Find a better quicker/less complex NN structure that converge and learn faster.
- Find better alpha and beta parameter that make PER more beneficial.
- Better understand the link between network structure and convergence.
- Understand why PER is bringing so much overfitting



On the vision problem:

- Change the input processing of the input image. Stop using RGB and create 1 frame for Yellow, one for Blue and one for background, on 4 frames, so 12 input channels to the convolutional network.
- Understand how long a network need to converge. (not like Deepmind brute force)
- Better understand the link between structure and convergence.
- Better understand impact/benefits of more depth or more convolutions.
- Better understand where to put residual layers.

I have spent some time on the code to make it nice, it is more now the fine tuning of parameters the next and perpetual challenge.